


```
In [1]: import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
from sklearn.datasets import load_iris
from sklearn.model_selection import train_test_split
```

Question 1.1

```
In [2]: a = 6
b = 4
c = 2
def f(x):
    return a + b*x + c*(x**2)
x = np.arange(-10, 10, 0.1)
y = np.vectorize(f)(x)
plt.plot(x,y)
```

Out[2]: [matplotlib.lines.Line2D at 0x10870abd490]



Question 1.2

```
In [3]: def grad_f(x):
    return 2*c*(x) + b
```

Question 1.3

The critical point is: $(-1, 4)$.

Question 1.4

```
In [4]: def grad_update(grad_func, x, eta):  
        return x - eta * grad_func(x)
```

Question 1.5:

```
In [5]: def min_find(x_init, epsilon, eta):  
        x_i = x_init  
        x_values = []  
        x_values.append(x_i)  
        x_i1 = grad_update(grad_f, x_i, eta)  
        while abs(x_i1 - x_i) > epsilon:  
            x_i = x_i1  
            x_values.append(x_i1)  
            x_i1 = grad_update(grad_f, x_i, eta)  
        return x_values  
  
x_final = min_find(100, 0.001, 0.01)[-1]  
print(f'({x_final}, {f(x_final)})')
```

```
(-0.9755852144560788, 4.001192163506311)
```

The x value is not the same as the one we found in question 1.3, but it is close. The reason for this being due to the epsilon being not zero.

Question 1.6

```
In [6]: T_min = np.inf  
T_min_x = None  
T_min_parameters = None  
  
# Loop over the range of epsilon, eta, and x_0  
for epsilon in np.arange(0.1, 0.0, -0.01):  
    for eta in np.arange(0.2, 0.0, -0.01):  
        for x_0 in np.arange(50, 5, -1):  
            ret = min_find(x_0, epsilon, eta)  
            if len(ret) < T_min:  
                T_min = len(ret)  
                T_min_x = ret  
                T_min_parameters = (x_0, epsilon, eta)  
  
# Print the results  
print(f'The algorithm is fastest with the following values:\n'  
      f'x_0 = {T_min_parameters[0]}\n'  
      f'epsilon = {T_min_parameters[1]}\n'  
      f'eta = {T_min_parameters[2]:.2f}\n'  
      f'With T = {T_min}')
```

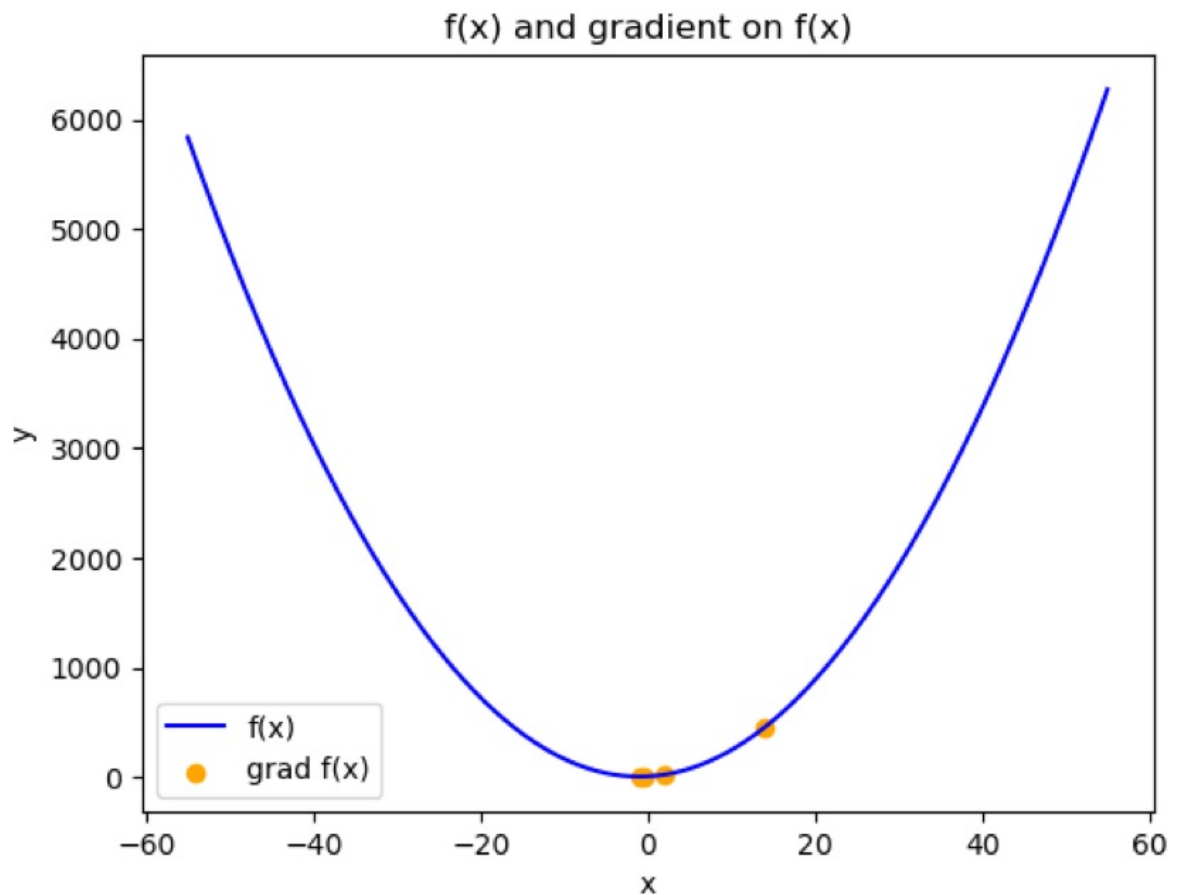
The algorithm is fastest with the following values:

```
x_0 = 14  
epsilon = 0.1  
eta = 0.20  
With T = 4
```

Question 1.7

```
In [7]: x = np.arange(-55,55,0.01)
y = np.vectorize(f)(x)
plt.plot(x,y,color='blue',label='f(x)')
y_min = np.vectorize(f)(T_min_x)
plt.scatter(T_min_x, y_min,color='orange',
            label='grad f(x)')
plt.title('f(x) and gradient on f(x)')
plt.xlabel('x')
plt.ylabel('y')
plt.legend()
```

Out[7]: <matplotlib.legend.Legend at 0x10870b0ba90>



1. ראשית נשים לב כי: $1 - y_i(\langle w, x_i \rangle + b)$ הינה פונקציה

אם קטורה. $\leq \max\{0, 1 - y_i(\langle w, x_i \rangle + b)\}$ קטורה

המקסימום של פונקציה קטורה היא קטורה גם כן.

$$\leq \frac{1}{m} \sum_{i=1}^m \max\{0, 1 - y_i(\langle w, x_i \rangle + b)\}$$

והכפלה בקבוע אי-שלילי תשאיר את התוצאה קטורה.

בנוסף $\lambda \|w\|^2 - \lambda$ קבוע אי-שלילי וטובה בהיקף

(קטורה כיוון $\|w\|$ קטורה ואי-שלילי מהצד השני) הינו קטורה

$$\leq \frac{1}{m} \sum_{i=1}^m \max\{0, 1 - y_i(\langle w, x_i \rangle + b)\} + \lambda \|w\|^2$$

זהו פונקציה קטורה

2. נסמן $f(w) = \ell(w, y_i, x_i)$

נחזיק:

$$\|f(w_1) - f(w_2)\| = \begin{cases} \|0 - 0\| \leq R \|w_1 - w_2\| & \forall R, f(w_1) = f(w_2) = 0 \\ \text{II} \quad \|1 - y_i \langle w_2, x_i \rangle\| \leq R \|w_1 - w_2\| & \text{if } f(w_1) = 0 \wedge f(w_2) > 0 \\ \text{III} \quad \|1 - y_i \langle w_1, x_i \rangle - (1 - y_i \langle w_2, x_i \rangle)\| \leq R \|w_1 - w_2\| & f(w_1), f(w_2) > 0 \end{cases}$$

$$\text{II} \Rightarrow 1 - y_i \langle w_1, x_i \rangle \leq 0 \Rightarrow 1 \leq y_i \langle w_1, x_i \rangle$$

$$\Rightarrow \|1 - y_i \langle w_2, x_i \rangle\| \leq \|y_i \langle w_1, x_i \rangle - y_i \langle w_2, x_i \rangle\|$$

$$\| \langle w_1 - w_2, y_i x_i \rangle \| \leq \|w_1 - w_2\| \cdot \max_i \|y_i x_i\| =$$

אם קושי-טורי

$$= \|w_1 - w_2\| \cdot \max_i \|x_i\|$$

בנוסף נחזיק: $(1 - y_i \langle w_2, x_i \rangle) > 0$ אם סדר III

$$\|1 - y_i \langle w_1, x_i \rangle - (1 - y_i \langle w_2, x_i \rangle)\| =$$

$$= \|y_i \langle w_1, x_i \rangle - y_i \langle w_2, x_i \rangle\| =$$

$$= \| \langle w_1 - w_2, y_i x_i \rangle \| \leq \|w_1 - w_2\| \cdot \max_i \|y_i x_i\| = \|w_1 - w_2\| \cdot \max_i \|x_i\|$$

אם קושי-טורי

Subgradient - ה , סימון λ 0.3

$$\text{subgradient by } w = \begin{cases} 2\lambda w & 1 - y_i(\langle w, x_i \rangle + b) \leq 0 \\ -y_i x_i + 2\lambda w & 1 - y_i(\langle w, x_i \rangle + b) > 0 \end{cases}$$

$$\text{Subgradient by } b = \begin{cases} 0 & 1 - y_i(\langle w, x_i \rangle + b) \leq 0 \\ -y_i & 1 - y_i(\langle w, x_i \rangle + b) > 0 \end{cases}$$

```
In [8]: def subgradient(w, x, y, b, lam):
        if 1 - y * (w.dot(x) + b) <= 0:
            return (2 * lam * w, 0)
        else:
            return (-y * x + 2 * lam * w, (-1) * y)
```

Question 2.4

```
In [46]: def svm_with_sgd(X, y, lam=0, epochs=1000, l_rate=0.01,
                        sgd_type='practical'):
    np.random.seed(2)
    m, d = X.shape
    b = np.random.normal(0, 1, 1)
    w = np.random.normal(0, 1, d)

    if sgd_type == 'practical':
        for epoch in range(epochs):
            for i in np.random.permutation(m):
                subgrad_w, subgrad_b =
                    subgradient(w, X[i], y[i], b, lam)
                w -= l_rate * subgrad_w
                b -= l_rate * subgrad_b
            return w, b
    if sgd_type == 'theory':
        ws, bs = [w.copy()], [b.copy()]
        # Use copy to avoid in-place modifications
        for _ in range(epochs * m):
            i = np.random.randint(m)
            subgrad = subgradient(w, X[i], y[i], b, lam)
            w -= l_rate * subgrad[0]
            ws.append(w.copy())
            b -= l_rate * subgrad[1]
            bs.append(b.copy())
        return np.mean(ws, axis=0), np.mean(bs, axis=0)
```

Question 2.5

```
In [38]: def calculate_error(w, bias, X, y):
    pred = []
    M = X.shape[0]

    for i in range(X.shape[0]):
        if w.dot(X[i]) + bias > 0:
            pred.append(1)
        else:
            pred.append(-1)
    return np.sum(pred != y) / M
```

Question 2.6

```
In [39]: # 2
X, y = load_iris(return_X_y=True)
X = X[y != 0]
y = y[y != 0]
y[y==2] = -1
X = X[:, 2:4]

# 3
X_train, X_val, y_train, y_val = train_test_split(X, y,
                                                    test_size=0.3, random_state=42)

# 4
lambdas = [0, 0.05, 0.1, 0.2, 0.5]
train_error_array = []
test_error_array = []
margins_array = []

for lam in lambdas:
    model = svm_with_sgd(X_train, y_train, lam)
    train_error = calculate_error(model[0], model[1],
                                   X_train, y_train)
    test_error = calculate_error(model[0], model[1],
                                   X_val, y_val)
    margin = 1 / np.linalg.norm(model[0])
    train_error_array.append(train_error)
    test_error_array.append(test_error)
    margins_array.append(margin)
```

```


In [40]: # 5
x = np.arange(len(lambdas))
width = 0.4

plt.bar(x, train_error_array, width=width,
        label='train', color='blue')
plt.bar(np.add(x, width), test_error_array,
        width=width, label='test', color='orange')

for i in range(len(lambdas)):
    plt.text(i, train_error_array[i] + 0.01,
             round(train_error_array[i], 2), ha='center', color='blue')
    plt.text(i + width, test_error_array[i] + 0.01,
             round(test_error_array[i], 2), ha='center', color='orange')

plt.xlabel('lambda')
plt.ylabel('error')
plt.title('2.6.a')
plt.xticks(np.add(x, width / 2), lambdas)
plt.legend()
plt.show()

```

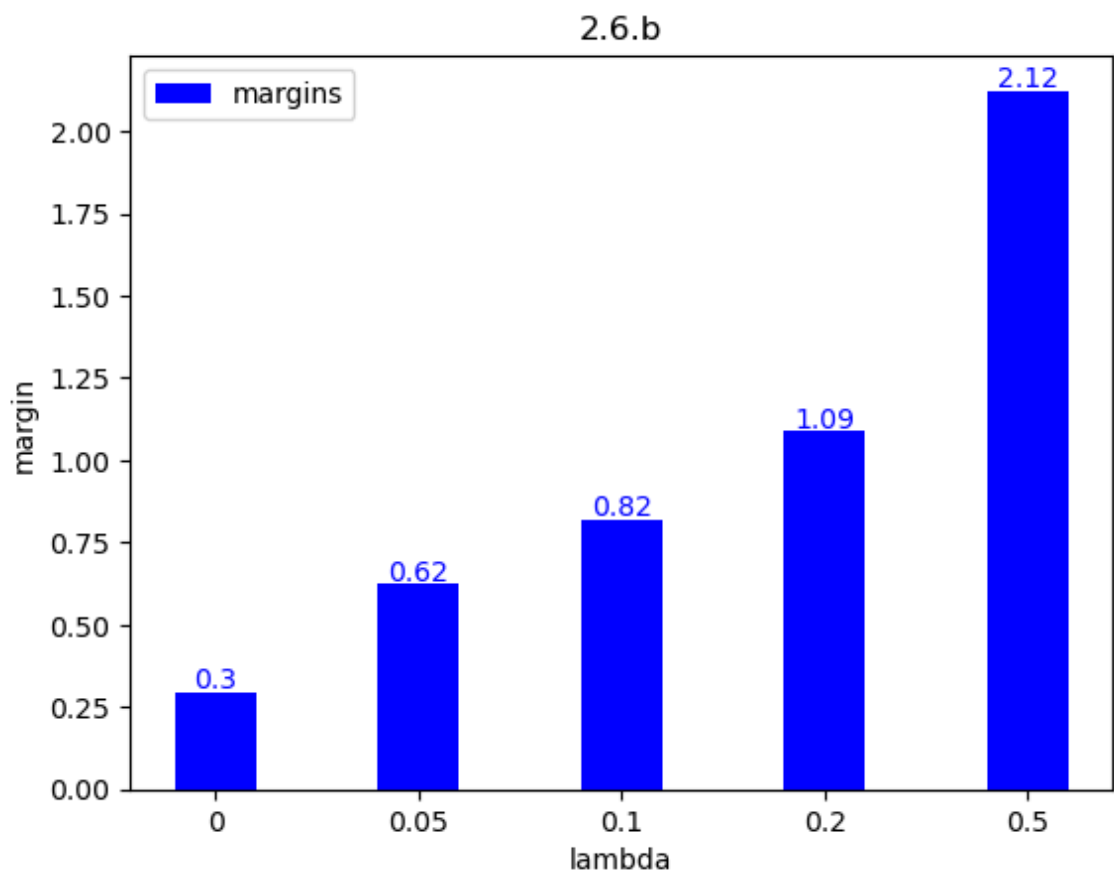


```
In [41]: x = np.arange(len(lambdas))
width = 0.4

plt.bar(x, margins_array, width=width, label='margins', color='blue')

for i in range(len(lambdas)):
    plt.text(i, margins_array[i] + 0.01,
             round(margins_array[i], 2), ha='center', color='blue')

plt.xlabel('lambda')
plt.ylabel('margin')
plt.title('2.6.b')
plt.xticks(x, lambdas)
plt.legend()
plt.show()
```



We choose $\lambda = 0.05$ since it is the λ with the lowest error rate for both the train and test sets.

Question 2.7

```
In [47]: epochs = np.arange(10, 1001, 10) # including 1000


train_error_practical = []
test_error_practical = []
train_error_theory = []
test_error_theory = []

for epoch in epochs:
    # Practical SGD
    model = svm_with_sgd(X_train, y_train,
                        lam=0.05, epochs=epoch, sgd_type='practical')
    train_error = calculate_error(model[0],
                                model[1], X_train, y_train)
    test_error = calculate_error(model[0],
                                model[1], X_val, y_val)
    train_error_practical.append(train_error)
    test_error_practical.append(test_error)

    # Theoretical SGD
    model = svm_with_sgd(X_train, y_train,
                        lam=0.05, epochs=epoch, sgd_type='theory')
    train_error = calculate_error(model[0],
                                model[1], X_train, y_train)
    test_error = calculate_error(model[0],
                                model[1], X_val, y_val)
    train_error_theory.append(train_error)
    test_error_theory.append(test_error)
```


```
In [48]: plt.plot(epochs, train_error_practical,  
                label='practical', color='blue')  
plt.plot(epochs, train_error_theory,  
         label='theory', color='orange')  
plt.xlabel('Epochs')  
plt.ylabel('Error')  
plt.title('2.7.a')  
plt.legend()
```

Out[48]: <matplotlib.legend.Legend at 0x10877805670>



```
In [49]: plt.plot(epochs, test_error_practical,
                 label='practical', color='g')
plt.plot(epochs, test_error_theory,
         label='theory', color='r')
plt.xlabel('Epochs')
plt.ylabel('Error')
plt.title('2.7.b')
plt.legend()
```

Out[49]: <matplotlib.legend.Legend at 0x108793d0cd0>



Theoretical SGD assumes noise-free, ideal data, leading to stable errors after a certain number of epochs. In contrast, practical SGD deals with real, noisy data, so errors can continue to change with more epochs due to inherent data imperfections. This explains why practical errors fluctuate while theoretical errors stabilize.

Question 3.1

```
In [17]: import numpy as np
```

```
In [18]: def cross_validation_error(X, y, model, n_folds):
    indices = np.arange(len(X))
    np.random.shuffle(indices)
    folds = np.array_split(indices, n_folds)

    train_errors = []
    validation_errors = []

    for fold in range(n_folds):
        val_indices = folds[fold]
        train_indices = np.concatenate([folds[i] for i in range(n_folds) if i != fold])

        X_train, y_train = X[train_indices], y[train_indices]
        X_val, y_val = X[val_indices], y[val_indices]

        model.fit(X_train, y_train)
        train_preds = model.predict(X_train)
        val_preds = model.predict(X_val)

        train_errors.append(np.mean(train_preds != y_train))
        validation_errors.append(np.mean(val_preds != y_val))

    avg_train_error = np.mean(train_errors)
    avg_val_error = np.mean(validation_errors)

    return avg_train_error, avg_val_error
```

Question 3.2

```
In [19]: import matplotlib.pyplot as plt
    from sklearn.svm import SVC
    from sklearn.datasets import load_iris
    from sklearn.model_selection import train_test_split
```

```
In [20]: # Define lambda values
lambda_values = [10**(-4), 10**(-2), 1, 10**2, 10**4]

def svm_results(X_train, y_train, X_test, y_test):
    results = {}
    n_folds = 5

    for lam in lambda_values:
        c_value = 1 / lam
        model = SVC(kernel='linear', C=c_value)

        train_error, val_error =
        cross_validation_error(X_train, y_train, model, n_folds)

        # Compute test error
        test_preds = model.predict(X_test)
        test_error = np.mean(test_preds != y_test)

        results[f'SVM_lambda_{lam}'] =
        (train_error, val_error, test_error)

    return results
```

Question 3.3


```
In [21]: # Load data
iris_data = load_iris()
X, y = iris_data['data'], iris_data['target']

# Split data
X_train, X_test, y_train, y_test = train_test_split(X, y,
                                                    test_size=0.2, random_s

# Evaluate models
results = svm_results(X_train, y_train, X_test, y_test)

# Plot results
x = np.arange(len(lambda_values))
train_errors = [results[f'SVM_lambda_{lam}'][0] for lam in lambda_values]
val_errors = [results[f'SVM_lambda_{lam}'][1] for lam in lambda_values]
test_errors = [results[f'SVM_lambda_{lam}'][2] for lam in lambda_values]

plt.figure(figsize=(10, 6))
plt.bar(x - 0.2, train_errors, width=0.2,
        label='Train Error', color='g')
plt.bar(x, val_errors, width=0.2,
        label='Validation Error', color='r')
plt.bar(x + 0.2, test_errors, width=0.2,
        label='Test Error', color='y')
plt.xlabel('Lambda')
plt.ylabel('Error Rate')
plt.title('SVM Model Errors for Different Lambda Values')
plt.xticks(x, lambda_values)
plt.legend()
plt.show()
```



The best model for the CV is $\lambda = 0.0001$. The best model for the test set is $\lambda = 1$. The λ which is best for CV is not the same as the one that is best for the test set, this is because for smaller λ s there is a significant difference between the train error and the test error - this is caused by overfitting. Since CV is based on the training set, the model that works best for it will be with smaller λ s. For a model which wishes to minimize the train error, $\lambda = 1$ is the most balanced between the smaller λ s (explained above) and larger λ s (which leads to underfitting).

$$g(w) + \langle u - w, \nabla g_j(w) \rangle = g_j(w) + \langle u - w, \nabla g_j(w) \rangle \leq g_j(u) \leq \max_i g_i(u) = g(u), \forall u \in \mathbb{R}^d$$

$g(w) = \max_i g_i(w) = g_j(w)$ מהצד

אם אי שווין שתיים במעלה
וההצדקה צדור פו קמורה

מהצדד $g_j(w)$ בשאלה